

# Sentiment Classification of Tweets About Products and Brands Using NLP





# Table of contents

**01**

**Introduction**

**02**

**Business  
Understanding**

**03**

**Problem  
Statement**

**04**


**Data  
Understanding**

**05**

**EDA, Findings  
Feature  
Engineering**

**06**

**Modelling**





# Business Understanding

In today's competitive tech landscape, brand perception plays a crucial role in product adoption, customer loyalty, and overall market performance. Social media platforms like Twitter provide real-time, unfiltered insights into how users feel about products and services. For global tech giants like **Apple** and **Google**, monitoring public sentiment around their products (e.g., iPhones, Android devices, iOS, Pixel phones, etc.) can help:

- Identify product issues early
- Understand customer preferences
- Measure the impact of product launches or controversies
- Inform marketing, PR, and customer support strategies

Given the vast volume and velocity of tweets, manual sentiment tracking is impractical — which makes automated sentiment classification an essential tool for brand managers and product teams.



# PROBLEM STATEMENT

The goal of this project is to develop a machine learning model that can automatically classify the sentiment of tweets related to Apple and Google products as Positive, Negative, or Neutral based on the content of the tweet.

## Stakeholders:


- Marketing teams
- Brand managers
- Customer experience teams





# Objectives

- **Develop a machine learning model** to classify whether a tweet contains an emotion directed at a product or brand, based on its text content.
- **Automated monitoring of consumer emotions** toward tech products (e.g., Apple, Google) on social media.
- **Perform Robust Data Preprocessing:** Implement thorough text cleaning, tokenization, stopword removal, and stemming to prepare high-quality textual data for modeling.
- **Engineer Predictive Features:**  
Extract relevant textual features (e.g., TF-IDF unigrams) and use mutual information for selecting the most informative terms.
- **Evaluate and Compare Models:**  
Train and evaluate multiple classification algorithms (e.g., Logistic Regression, Random Forest) and select the best-performing one based on accuracy, precision, recall, and F1-score.



# Data Understanding



Link : <https://data.world/crowdfunder/brands-and-product-emotions>


The dataset contains over 9,000 tweets mentioning Apple or Google products, labeled with whether an emotion is present and directed at a brand or product. Each tweet includes text, a target entity (like iPhone or Google), and a sentiment label such as "Positive emotion," "Negative emotion," or "No emotion."

## Features Before Cleaning

1. **tweet\_text**: the raw tweet
2. **emotion\_in\_tweet\_is\_directed\_at**: specific product/brand (e.g., iPhone, Google)
3. **Is\_there\_an\_emotion\_directed\_at\_a\_brand\_or\_product**: sentiment label

## Final Features After Cleaning


1. **text** – original tweet text
2. **target** – specific product/brand
3. **sentiment** – cleaned label (3-class)
4. **brand** – mapped brand (Apple, Google, None)
5. **tweet\_length** – number of characters in the tweet
6. **word\_count** – number of words in the tweet  
**clean\_text** – lowercase punctuation/URL/mention-free text
7. **tokens** – tokenized version of clean text
8. **preprocessed\_text** – tokenized + stopwords removed + stemmed





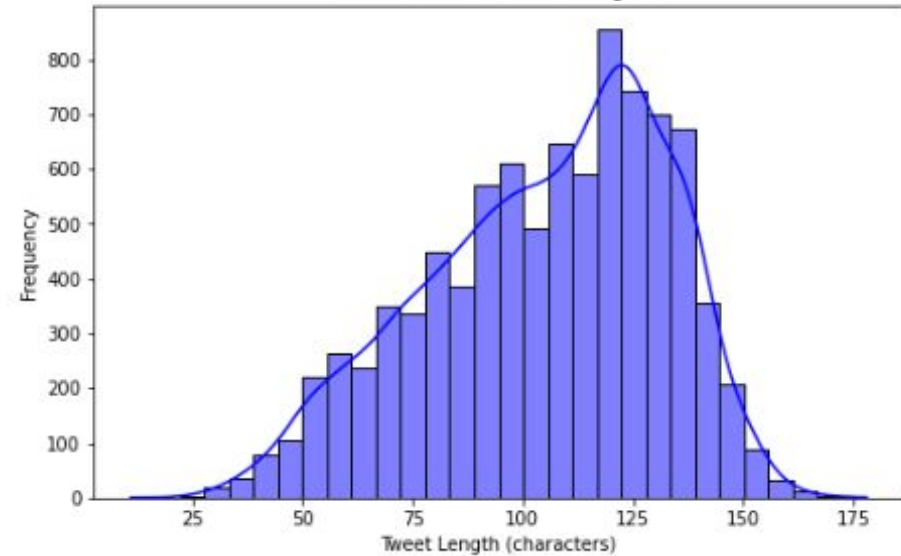
# Data Preprocessing Steps

- Lowercasing- All text was converted to lowercase to ensure consistency in word representation
- Removing Punctuation, Mentions, URLs
- Whitespace Normalization- Extra spaces were cleaned up using regex to make text uniform.
- Tokenization- Cleaned text was split into individual words (tokens), forming a list of tokens for each tweet.
- Stopword Removal- Common English stopwords (e.g., "the", "is", "and") were removed using NLTK's stopwords list to reduce noise.
- Stemming- Words were reduced to their root forms using PorterStemmer (e.g., "running" → "run", "tweets" → "tweet").
- Creating Final Preprocessed Text- The cleaned, stemmed tokens were joined back into a single string (preprocessed\_text) to be used in vectorization.
- TF-IDF Vectorization- Used TfidfVectorizer with unigrams (single words) to convert text into numerical features for model input.
- Train-Test Split- Final dataset was split into training and testing sets (80/20) for modeling.

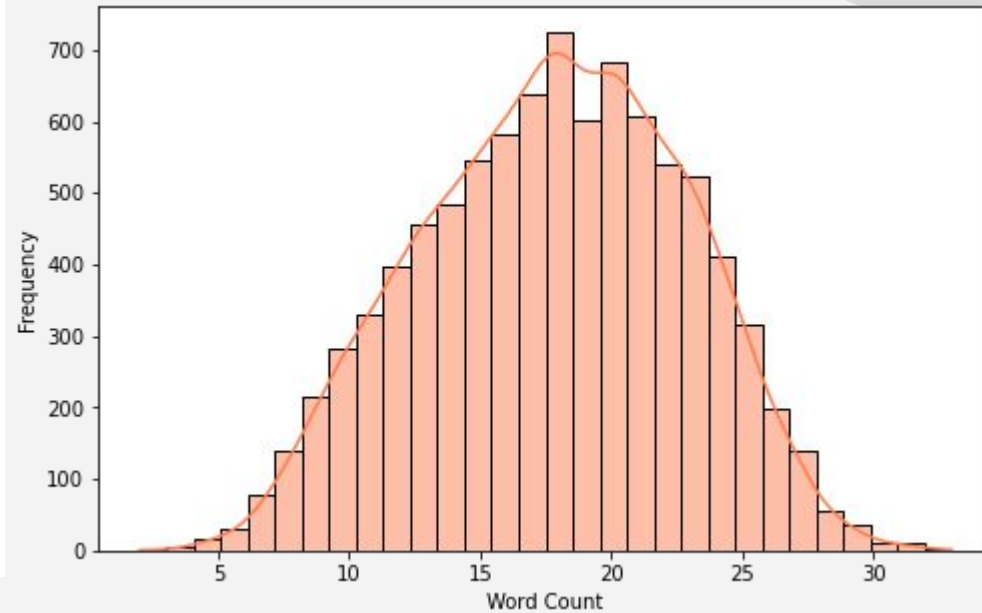


# EXPLORATORY DATA ANALYSIS

Distribution of Tweet Lengths




Distribution of Word Counts in Tweets





# Word Count by Sentiment

Emotional tweets tend to have more words, suggesting wordiness correlates with sentiment.





# Data Preprocessing Steps

## Modeling Approach

### 1. Problem Framing

- Binary Classification: Predict positive (1) vs. negative (0) sentiment.
- Multiclass Classification: Extend to 3 classes (positive, negative, neutral).
- Class Imbalance Addressed: Used `class_weight='balanced'` and oversampling for fairness.

### 2. Models Evaluated

#### Binary Classification:

- Baselines:
  - Naive Bayes (poor negative recall: 2%)
  - Logistic Regression (untuned: 11% negative recall; tuned: 61%)
  - Random Forest (25% negative recall)
  - SVM (best baseline: 58% negative recall, 88% accuracy)






# Data Preprocessing Steps

- Neural Network:
  - Architecture: Dense(128) → Dropout(0.3) → Dense(64) → Dropout(0.3) → Sigmoid
  - Result: Severe overfitting (100% train/test accuracy).



Multiclass Classification (positive, negative, neutral):

- Neural Network:
  - Architecture: Dense(64, L2) → Dropout(0.5) → Dense(32, L2) → Dropout(0.5) → Softmax
  - Regularization: L2 penalty + dropout to combat overfitting.
  - Result: 100% test accuracy (potential overfitting despite regularization).





# FLOWCHART





**THANK YOU  
ANY  
QUESTIONS?**

